Gur Raunaq Singh

+91-9015154091 | raunaq.soni@gmail.com | linkedin.com/in/gurraunaqsingh | github.com/raunaqness | Bangalore, India

Summary

Senior AI and ML Systems Architect with 8+ years of experience designing secure, scalable AI platforms in enterprise environments processing high-volume transactional data. Specialized in distributed ML systems, Kubernetes-native orchestration, multi-cloud architectures, and Generative AI (LLMs, RAG, agentic workflows). Strong focus on MLOps, IAM-based access control, auditability, and operational resilience. Experience deploying and managing containerized workflow automation systems (n8n) on cloud infrastructure with secure API and LLM integrations.

Technical Skills

AI/ML: LLMs, RAG, LangChain, LangGraph, Embeddings, Semantic Search, Hybrid Retrieval, Reranking, Prompt Engineering, Model Monitoring

Cloud & Infra: Azure (AKS, App Services), AWS, Kubernetes, Docker, Terraform

Data & Backend: Python, SQL, Kafka, Redis, PostgreSQL, FastAPI, Django, API Development, Webhooks

Vector Databases & Search: Pinecone DB, Vector Search, Metadata Filtering, Document Chunking, Retrieval Pipelines

MLOps: CI/CD, MLflow, IAM, Audit Logging, Encryption, SLA Monitoring

Automation: n8n (workflow automation), API integrations, Webhooks

Experience

Network School — Jun 2026 – Present
Resident — *Singapore (Remote)*

- Working as a consultant helping teams design and build AI systems, RAG systems, and multi-agent systems.
- Designing and developing AI-powered tools, RAG-based chatbots, and AI agents using Python, LangChain, LangGraph, and FastAPI, with a focus on rapid prototyping and scalable deployment.
- Building and documenting an evolving portfolio of hands-on projects, technical explorations, and open-source contributions: raunaqness.com/projects

Aptos India — Jan 2023 – Present
Machine Learning Engineer — *Bangalore, India*

Project: Aptos One – Enterprise Unified Commerce Platform serving 100+ global retailers and processing billions of transactions annually.

- Architected a Kubernetes-native ML orchestration control plane enabling workload-aware scheduling, checkpoint-based recovery, and distributed job lifecycle management in SLA-bound production environments.
- Designed secure, fault-tolerant data pipelines on GCP (BigQuery + event-driven workflows), implementing IAM-based access controls, encryption practices, and audit logging while improving processing SLAs by 35%.
- Optimized cloud infrastructure and query execution strategies, reducing compute costs by 60% (about USD 240K annual savings) while modernizing legacy systems into scalable, observable cloud-native components.
- Technologies: Python, GCP, Vertex AI, BigQuery, Kubernetes, Docker, Terraform, Redis, CI/CD, IAM

Project: Generative AI & LLM Support Assistant

- Designed an enterprise-grade RAG architecture integrating structured analytics metadata and internal knowledge bases to deliver context-aware responses with improved semantic accuracy.
- Implemented prompt versioning, evaluation workflows, and monitoring mechanisms to reduce hallucinations by 35% while improving traceability and observability of model outputs.
- Deployed containerized inference workflows using MLOps practices, enabling controlled experimentation, scalable inference, and secure knowledge isolation via vector database access controls.
- Technologies: Python, LangChain, LangGraph, Vertex AI, VectorDB, FastAPI, Docker, MLOps

Turing.com — Apr 2022 – Dec 2022
Senior Software Engineer — *India (Remote)*

Project: End-to-End ML Deployment Platform

- Designed and implemented production-grade ML pipelines using Airflow for ingestion, training, validation, and automated deployment.
- Standardized CI/CD and model versioning practices, reducing deployment time by 60% and improving controlled production rollouts.
- Integrated monitoring and validation checks to enhance production reliability and maintainability of deployed ML systems.
- Technologies: Python, Django, Airflow, GCP, Docker, MLOps

Clickpost Pvt. Ltd. Apr 2021 – Apr 2022
Senior Software Engineer *Bangalore, India*

Project: Cash Reconciliation Automation Platform – serving 1,000+ enterprise customers handling high-volume CoD transactions.

- Designed and built a high-accuracy financial reconciliation system processing large-scale transactional datasets with strict data integrity and traceability requirements.
- Reduced reconciliation cycle time by 95% (from about 2 days to minutes) through parallelized processing and automated validation workflows, improving operational efficiency and financial visibility.
- Architected secure ingestion pipelines and scalable backend APIs under high-availability constraints, ensuring auditability and controlled access to financial data.
- Technologies: Python, Django, PostgreSQL, AWS, Redshift, Kafka, Docker, Kubernetes

Veda Labs Pvt. Ltd. Jul 2018 – Apr 2021
Software Engineer *New Delhi, India*

Project: Video Analytics Platform – Computer vision platform deployed on-premise for retail analytics using CCTV footage.

- Developed containerized video analytics pipelines for edge devices.
- Built a Kafka-based streaming architecture for scalable, real-time processing.
- Technologies: Python, OpenCV, Computer Vision, Kafka, Docker

Projects

FluxBox – macOS Productivity Menu Bar App (Rust + Tauri)

- Ultra-lightweight macOS menu bar utility with real-time market tickers, currency conversion, world clocks, hardware monitoring, and Claude API usage tracking. Built with Rust and Tauri for near-zero resource footprint.
- Github: https://github.com/raunaqness/fluxbox

Stateful AI Travel Agent (LangGraph)

- Engineered a graph-orchestrated, memory-aware multi-agent system for dynamic travel planning.
- Github: https://github.com/raunaqness/ai-travel-agent
- Demo: https://ai-travel-agent.fly.dev/

Multimodal RAG Platform

- Designed and built an enterprise-grade Retrieval-Augmented Generation (RAG) system for unstructured document search and question answering, delivering grounded, context-aware responses with reduced hallucinations.
- Implemented end-to-end retrieval pipelines including document ingestion, chunking, embedding generation, vector indexing in Pinecone DB, semantic search, metadata filtering, and reranking for high retrieval accuracy.
- Orchestrated LLM prompt workflows with citation-backed responses, observability, evaluation, and secure multi-tenant access controls, deploying the solution as a scalable production-ready API.
- Github: github.com/raunaqness/rag | Demo: rag.raunaqness.com

Certifications

- Building with the Claude API – Anthropic verify.skilljar.com/c/ufyhuqnu7z75
- Claude Code in Action – Anthropic verify.skilljar.com/c/5wq8kmpa2jbk
- Introduction to Agent Skills – Anthropic verify.skilljar.com/c/f5vktmpg7und
- Introduction to Model Context Protocol – Anthropic verify.skilljar.com/c/73xymxvikxt7

Education

Indraprastha University Aug 2014 – Jun 2018
Bachelor of Technology – Computer Science Engineering *New Delhi, India*